

FINM 32950: Intro to HPC in Finance

Lecture 3

Chanaka Liyanarachchi

April 4, 2025

Structured Parallel Programming

High Level Constructs for Parallel Programming

Monte Carlo

Structured Parallel Programming

Road Map

We discussed some important concepts about multithreading:

- ▶ creating threads
- ▶ thread synchronization
- ▶ critical sections
- ▶ race conditions
- ▶ mutual exclusions
- ▶ locks
- ▶ deadlocks
- ▶ atomics

Problem

- ▶ Our goal: use threads to write parallel programs.
- ▶ Suppose we profile an application and see low CPU utilization:

Basic Hotspots Hotspots by CPU Usage viewpoint (change) ⓘ

Analysis Target Analysis Type Collection Log Summary Bottom-up Caller/Callee Top-down Tree Pl

Elapsed Time ⓘ: 1.575s

CPU Time ⓘ: 1.003s
Total Thread Count: 1
Paused Time ⓘ: 0s


Top Hotspots

This section lists the most active functions in your application. Optimizing these hotspot functions typically results in improving ov

Function	Module	CPU Time ⓘ
rand	ucrtbased.dll	0.336s
cdf_normal	BlackScholes.exe	0.120s
logf	BlackScholes.exe	0.106s
__libm_exp_s9	libmmd.dll	0.090s
__libm_log_s9	libmmd.dll	0.080s
[Others]		0.270s

CPU Usage Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhe



- ▶ Which part/parts can be run in parallel?


Pattern Based Approach

- ▶ We will learn how to do this in a structured way using a *pattern-based* approach.
- ▶ Patterns are widely used in software engineering.
- ▶ Capture best practices to solve common problems in a particular domain.
- ▶ *Parallel Patterns*: capture known solutions and best practices to solve known problems in parallel program design.
- ▶ We can use one or more parallel patterns to design a program.
- ▶ In this course, we will:
 - ▶ introduce some common parallel patterns
 - ▶ implement them (homework)
 - ▶ use them as high level design tools to write parallel programs

Parallel Patterns

Folk-join

- ▶ Very common programming pattern as shown below¹




- ▶ Generally, this pattern involves:
 - ▶ serial work
 - ▶ parallel work
 - ▶ wait for all parallel work to finish
 - ▶ go to next task


¹figure source: wikipedia

Map

- ▶ Map pattern involves applying an operation to all elements in a sequence.
- ▶ This operation produces a new collection with the same shape as the input.
- ▶ Serial execution of this pattern takes the form of:



- ▶ Parallel execution of this pattern takes the form of:




Map

- ▶ The operation is known as *elemental function*.
- ▶ The elemental function should not modify shared (global) data that other instances depend on.
- ▶ E.g.: Assignment A, Assignment B

Stencil

- ▶ Here the elemental function access an element in a sequence and its neighbors.



- ▶ E.g.: Moving average, finite difference.

Reduction

- ▶ A reduction combines elements in a collection into a single element.
- ▶ A simple example would be computing the sum of elements in a vector.
- ▶ A function used in reduction is known as the *combiner function*.
- ▶ E.g.: Monte Carlo.


Pipeline

- ▶ Pipeline is a very common example/concept in everyday life.
- ▶ E.g. Assembly line of a car factory.



- ▶ Performs a specific task on an item at each station.
- ▶ Move the item to the next station for the next task.

- ▶ Next item takes the place of an item when it leaves a station.
- ▶ All stations are busy.



- ▶ Widely used in Processor/FPGA designs.

Many Patterns

- ▶ We looked at some patterns we plan to use in this course.
- ▶ There are many more:
 - ▶ Scan
 - ▶ Scatter
 - ▶ Gather
 - ▶ Futures
 - ▶ Search
 - ▶ ...
- ▶ A detailed discussion of all common parallel patterns is beyond the scope of this course.
- ▶ Recommended reading: *Structured Parallel Programming*² is a great resource for further studies on the subject.

²See course references for details

Assignment 1 - Required

- ▶ Due: April 12 by midnight.
- ▶ Implement any 2 patterns below using threads (`std::thread`) to demonstrate how the pattern works.
 - ▶ Folk-join
 - ▶ Map
 - ▶ Stencil
 - ▶ Reduction
- ▶ Your program doesn't have to do anything *meaningful*. Great if it does.
- ▶ For example, an elemental function can just be a `cout` statement.

Using Threads: Difficulties

- ▶ We have seen how to use threads to do work in parallel.
- ▶ Using `std::thread` is not easy (low level construct):
 - ▶ How many threads to use is not clear
 - ▶ Load balancing difficult
 - ▶ Doesn't scale
- ▶ Topics we discussed are very important:
 - ▶ Need a good foundation to understand how threads are used to write parallel programs
 - ▶ Learn related concepts: critical sections, data races, locks etc.
- ▶ We are going to discuss a different way/approach to use threads next.

High Level Constructs for Parallel Programming

High Level Constructs for Concurrency

- ▶ We will look at some *high level programming constructs* for parallel programming.
- ▶ Express parallelism at high-level hiding low-level implementation details.
- ▶ Useful for the beginner and the experienced programmer, alike.
- ▶ Advantages:
 - ▶ Clarity
 - ▶ Ease of use (increase productivity)
 - ▶ Less error prone
 - ▶ Scalability

- ▶ Popular high-level constructs:
 1. OpenMP: very popular
 2. (Intel) Threading Building Blocks (TBB): many features
 3. (Microsoft) Parallel Patterns Library (PPL): features similar to TBB
 4. C++ Standard Library
- ▶ We don't have to use only one solution. We can mix-and-match.

OpenMP (**O**pen specification for **M**ulti **P**rocessing)

OpenMP

- ▶ Very popular model.
- ▶ Provides a simple approach to develop parallel code, using:
 1. compiler directives
 2. functions
 3. environment variables
- ▶ We don't have to use all 3 of them (some features overlap).
- ▶ Implemented based on a specification³ developed by industry participants, including, IBM, Intel, Oracle and several others⁴.

³www.openmp.org

⁴Microsoft is no longer a member.

- ▶ Write parallel code in a platform (compiler, OS) neutral manner:
 - ▶ Works on many OSs:
 1. Windows
 2. (Various flavors of) Unix and Linux
 - ▶ Popular compiler support:
 1. Visual C++ supports OpenMP 2.0⁵
 2. <https://devblogs.microsoft.com/cppblog/improved-openmp-support-for-cpp-in-visual-studio/>
 3. Intel, gcc are more up-to-date.
 - ▶ Supports different programming languages:
 1. C++
 2. C
 3. Fortran

⁵<https://msdn.microsoft.com/en-us/library/0ca2w8dk.aspx>

OpenMP: Example 1

- ▶ Code below shows a simple *Hello World* program written using OpenMP directives:

```
#include <omp.h>

int main()
{
    #pragma omp parallel
    {
        cout << "Hello, world" << endl;
    }
}
```


- ▶ The OpenMP uses compiler directives to make regions of code execute in parallel:
 - ▶ We don't create threads explicitly
 - ▶ We instruct the compiler what to do
 - ▶ Compiler generates appropriate code
- ▶ OpenMP features are defined in `<omp.h>`.

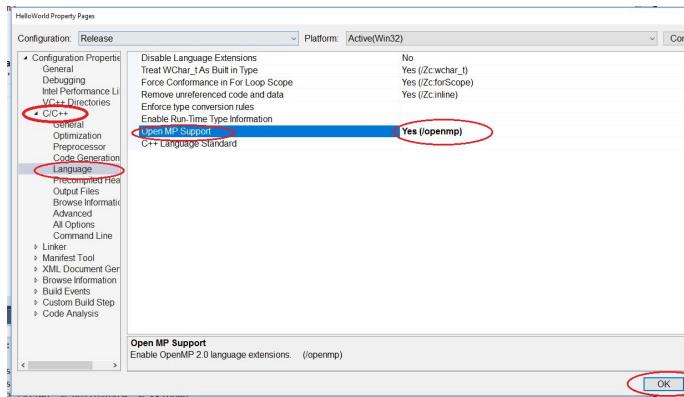
Enabling OpenMP Support: Midway

- ▶ Intel Compiler (Linux):
- ▶ use -qopenmp flag:
`icc -qopenmp hello.cpp -o hello`
- ▶ Note: You need to load intel module first (on midway).

Enabling OpenMP Support: Windows

- ▶ We also need to enable OpenMP support.
- ▶ Visual C++ (Windows):

Properties -> C/C++ -> Language -> OpenMP Support -> Yes



OpenMP Directive Format

An OpenMP directive follows the format:

```
#pragma omp directive-name [clause]
```

- ▶ `#pragma` provides a mechanism to provide additional information to the compiler.
- ▶ `#pragma omp`: tells the compiler it is an OpenMP directive. This field is required.
- ▶ `directive-name`: a valid OpenMP directive must appear after the pragma. This field is required.
- ▶ `[clauses]`: tells the more about what has to be done. This field is optional.

OpenMP: Example 2

- For example, we can use an optional clause to specify the number of threads:

```
#include <omp.h>

int main()
{
    #pragma omp parallel num_threads(2)
    {
        cout << "Hello, world" << endl;
    }
}
```

OpenMP: Benefits

- ▶ OpenMP offers many features to write parallel code without using low level constructs.
- ▶ We still need to pay attention to:
 - ▶ Race conditions – when shared data members are accessible (read and write) by more than one thread
 - ▶ Dependencies – multiple threads in a block may run in any order
- ▶ References:
 - ▶ <https://www.openmp.org/>
 - ▶ Quick reference: <http://www.openmp.org/wp-content/uploads/OpenMP-4.0-C.pdf>
 - ▶ Text: *Using OpenMP*⁶

⁶See course references list for details

OpenMP Directives

Parallel:

- ▶ The `parallel` directive defines a parallel region:
 - ▶ At the beginning of the `parallel` directive block, worker threads are created
 - ▶ Each thread executes every instruction in the parallel block
 - ▶ At the end of the parallel block, all threads join (i.e. there is an implicit barrier).
- ▶ The order in which the threads run is not specified.
- ▶ To provide an explicit barrier:
`#pragma omp barrier`

Parallel Loops:

- ▶ `parallel for`:
 - ▶ used to parallelize a for loop
- ```
#pragma omp parallel for
```

- ▶ Example:

```
#pragma omp parallel for
for (int i=0; i<10; ++i)
{
 c[i] = a[i] + b[i];
}
```



## OpenMP: Example 3

- ▶ Using OpenMP, multithreading Matrix multiplication is very easy. Thread-per-row can be implemented easily as follows:

```
void matrix_multiply_parallel(const matrix& m1,
 const matrix& m2,
 matrix& m3,
 int rows, int columns)
{
 #pragma omp parallel for
 for (int i = 0; i < rows; ++i)
 {
 for (int j = 0; j < columns; ++j)
 {
 m3[i][j] = 0;
 for (int k = 0; k < rows; ++k)
 {
 m3[i][j] += m1[i][k] * m2[k][j];
 }
 }
 }
}
```

- ▶ Exercise: Change the number of threads; observe execution times.

# Sections

- ▶ Parallel for allows us to execute the same task on different data in parallel easily.
- ▶ We use sections to execute different tasks (different kinds of work) in parallel:
  - ▶ Specifies different code regions.
  - ▶ Each region is executed by one thread.
  - ▶ Each section should run independently of the other regions.

► Example:

```
#pragma omp parallel
{
 #pragma omp sections
 {
 #pragma omp section
 task1();

 #pragma omp section
 task2();
 }
}
```

- Here, task1 and task2 will be executed by two threads in parallel.

# Thread Scheduling/Load Balancing

- ▶ OpenMP supports two thread scheduling policies:
  - ▶ static
  - ▶ dynamic
- ▶ Helpful in *load balancing*.
- ▶ Static scheduling (default):  
`#pragma omp parallel for schedule(static)`
  - ▶ Preallocates threads to loop indexes.
  - ▶ If one thread finishes first, the loop indexes won't be re-assigned.
- ▶ Dynamic scheduling:  
`#pragma omp parallel for schedule(dynamic)`
  - ▶ Allocates loop indexes to threads as threads become available to perform work.

- ▶ Loop below shows some uneven work.
- ▶ Let's use static scheduling first:

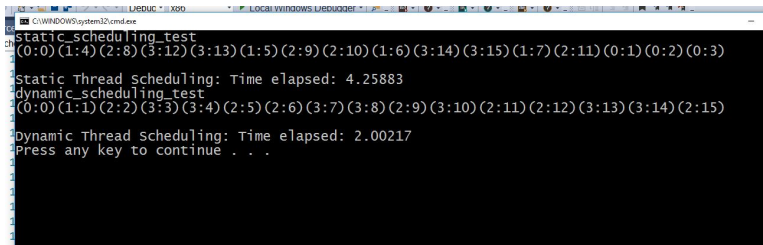
```
#pragma omp parallel for schedule(static) num_threads(4)
for (int i = 0; i<16; ++i)
{
 if (i<2)
 {
 Sleep(2000); // simulate long work
 }
 else
 {
 Sleep(100); // simulate short work
 }
 #pragma omp critical
 cout << "(" << omp_get_thread_num()
 << ":" << i << ")" << flush;
}
```

- ▶ A global lock (omp critical) is used to write to cout in an orderly fashion.

- And, dynamic scheduling next:

```
#pragma omp parallel for schedule(dynamic) num_threads(4)
for (int i = 0; i<16; ++i)
{
 if (i<2)
 {
 Sleep(2000); // simulate long work
 }
 else
 {
 Sleep(100); // simulate short work
 }
 #pragma omp critical
 cout << "(" << omp_get_thread_num()
 << ":" << i << ")" << flush;
}
```

- ▶ Static scheduling:
  - ▶ Each thread completes 4 tasks each
  - ▶ Number of work items equally shared
- ▶ Dynamic scheduling:
  - ▶ Allocates loop indexes to threads as threads become available to perform work
  - ▶ Some threads do more work items than others



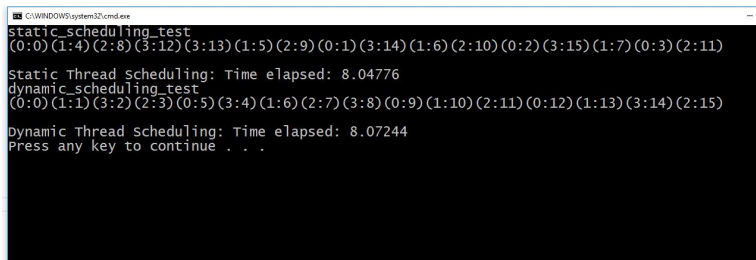
```
C:\WINDOWS\system32\cmd.exe
static_scheduling_test
(0:0) (1:4) (2:8) (3:12) (3:13) (1:5) (2:9) (2:10) (1:6) (3:14) (3:15) (1:7) (2:11) (0:1) (0:2) (0:3)
Static Thread Scheduling: Time elapsed: 4.25883
dynamic_scheduling_test
(0:0) (1:1) (2:2) (3:3) (3:4) (2:5) (2:6) (3:7) (3:8) (2:9) (3:10) (2:11) (2:12) (3:13) (3:14) (2:15)
Dynamic Thread Scheduling: Time elapsed: 2.00217
Press any key to continue . . .
```

- ▶ Dynamic scheduling works better for uneven work.

- ▶ If we have similar work, static scheduling works better:

```
#pragma omp parallel for schedule(static) num_threads(4)
for (int i = 0; i<16; ++i)
{
 Sleep(2000); // simulate work

 #pragma omp critical(log)
 cout << "(" << omp_get_thread_num()
 << ":" << i << ")" << flush;
}
```



```
static_scheduling_test
(0:0)(1:4)(2:8)(3:12)(3:13)(1:5)(2:9)(0:1)(3:14)(1:6)(2:10)(0:2)(3:15)(1:7)(0:3)(2:11)

Static Thread Scheduling: Time elapsed: 8.04776
dynamic_scheduling_test
(0:0)(1:1)(3:2)(2:3)(0:5)(3:4)(1:6)(2:7)(3:8)(0:9)(1:10)(2:11)(0:12)(1:13)(3:14)(2:15)

Dynamic Thread Scheduling: Time elapsed: 8.07244
Press any key to continue . . .
```

- ▶ We're using a named lock (`omp critical(log)`) here to write to `cout` in an orderly fashion.



# Synchronization Constructs

- ▶ `barrier`:
  - ▶ Provides a point of synchronization for threads in the group.
- ▶ `critical`:
  - ▶ A critical section without a name acts as a global lock
- ▶ `critical(lock_name)`:
  - ▶ A named critical section serializes execution with other critical sections with the same name
  - ▶ Critical sections with different names can run concurrently
- ▶ `master`:
  - ▶ This code block is executed by the master thread only
- ▶ `atomic`: Next slide

- ▶ An operation to be executed atomically (all or none, without interruptions).
- ▶ Operation supports additional clauses: read, write, update and capture.
- ▶ Counter example (see demo for details):

```
#pragma omp parallel for
for (int i = 0; i < 100000; ++i)
{
 #pragma omp atomic
 counter++;
}
```

## Reductions

- ▶ Shown below is a simple example for reduction:

```
sum = 0;
for (int i = 0; i < 10; i++)
{
 sum += a[i];
}
```

- ▶ sum is a shared variable – doing this in parallel is not trivial.
- ▶ The OpenMP reduction clause is used to accumulate a value in parallel easily.

```
int sum_sq = 0;

#pragma omp parallel for reduction(+:sum_sq)
for (int i = 0; i < 5; ++i)
{
 sum_sq += i*i;
}

std::cout << sum_sq << std::endl;
```

- ▶ Reduction is supported for following operators:  
+, -, \*, &, |, ^, &&, ||

## Multithreading Examples

# Matrix Multiplication: Serial to Parallel

- ▶ We matrix multiplication to illustrate parallel programming models.
- ▶ Serial:

```
void MatrixMultiply(const matrix& m1, const matrix& m2,
 matrix& m3,
 int rows, int columns)
{
 for (int i = 0; i < rows; ++i)
 {
 CalculateRow(/* args */);
 }
}
```

► Using `std::thread`:

```
void MatrixMultiply(const matrix& m1, const matrix& m2,
 matrix& m3,
 int rows, int columns)
{
 vector<thread> threads(rows);

 for (int i = 0; i < rows; ++i)
 {
 threads[i] = thread([&]()
 {
 CalculateRow(/* args */);
 });
 }

 for (thread& t : threads)
 {
 t.join();
 }
}
```

► Using OpenMP:

```
void MatrixMultiply(const matrix& m1, const matrix& m2,
 matrix& m3,
 int rows, int columns)
{
 #pragma omp parallel for
 for (int i = 0; i < rows; ++i)
 {
 CalculateRow(/* args */);
 }
}
```

► OpenMP solution has following advantages:

- Clarity
- Ease of use (increase productivity)
- Less error prone
- Scalability

## Intel Parallel Programming Models



- ▶ Intel toolkit offers a wide range of support for high-performance and parallel computing.
- ▶ We've seen some features so far:
  - ▶ Tools: e.g. vtune for profiling, advisor for vectorization help
  - ▶ SVMML
- ▶ Now, we will introduce two useful features in the Intel toolkit and illustrate how to use them in parallel Monte Carlo simulations:
  - ▶ Threading Building Blocks (TBB)
  - ▶ Math Kernel Library (MKL)

## Threading Building Blocks (TBB)

# TBB

- ▶ A library solution from Intel for high level *task based* parallel programming .
- ▶ Supports a wide range of features useful for parallel programming.
- ▶ Developer guide: <https://www.intel.com/content/www/us/en/docs/onetbb/developer-guide-api-reference/2021-6/onetbb-developer-guide.html>

# Task Based Programming

- ▶ The task\_group class can be used to run tasks in parallel.

```
tbb::task_group g;
```

```
g.run([&]() { work1(); });
```

```
g.run([&]() { work2(); });
```

```
g.run([&]() { work3(); });
```

```
g.wait();
```

## Detour: Function Objects (C++)

# Function Objects

- ▶ Function objects are objects that behave like functions.
- ▶ We get this behavior when we overload `operator()` for the objects of a class.
- ▶ Function objects can have other member functions and member variables – they are functions with state.
- ▶ Function objects are used in various places in C++.
- ▶ We're going to use them with `threads/TBB`.

## Creating a Thread

- ▶ We can create a thread using a function object:

```
class ThreadTask
{
private:
 int threadID_;

public:
 ThreadTask(int id)
 : threadID_(id)
 {}

 void operator()() const
 {
 std::cout << "Hello, World, " << threadID_ << std::endl;
 }
};
```

- ▶ Usage:

```
std::thread t(ThreadTask(2));
t.join();
```

- ▶ We're passing an object to thread constructor.
- ▶ Thread executes the overloaded `operator()`.

# Task Based Programming

- ▶ With TBB tasks:

```
const int NumTasks = 10;
```

```
tbb::task_group group;
```

```
for (int i = 0; i < NumTasks; ++i)
{
 group.run(ThreadTask(i));
}
```

```
group.wait();
```



# Task Based Programming: Advantages

- ▶ Tasks allow us to design a program using logical (high level) tasks.
- ▶ As we saw earlier, the number of threads we use is important:
  - ▶ Create too many – poor performance due to overhead
  - ▶ Don't create enough – wasting cpu resources
  - ▶ Ideal number depends on the system
- ▶ When we use tasks, we don't need to figure out the exact number of threads.

## Task scheduler:

- ▶ Creates the correct number of threads based on system resources
- ▶ Does load balancing
- ▶ Manages threads internally:
  - ▶ Doesn't create a thread every time we want to run task
  - ▶ Less thread creation overhead (i.e. creating a thread can be expensive compared to tasks)
- ▶ Schedules threads more efficiently:
  - ▶ Threads – OS uses fair scheduling; can be less efficient for some cases
  - ▶ Task scheduler claims to schedule threads more “intelligently” (using unfair scheduling). e.g. doesn't start a task until it can make progress

Monte Carlo

- We know how to write a serial version of Monte Carlo Option pricer to price European call options:

```
for (unsigned int i = 0; i < M; ++i)
{
 double z_i = generate_random_number();

 double ST_i = S0*exp((r - sigma*sigma / 2.0)*T
 + sigma*z_i*sqrt(T));

 sumPayoffs += (ST_i > K) ? ST_i - K: 0.0;
}

return exp(-r*T)*(sumPayoffs / M);
```

# Generating Random Numbers

- ▶ We have many options to generate pseudo random numbers (discussed in Computing for Finance):
  - ▶ Implement an algorithm, e.g. Box Muller
  - ▶ Use generators in the C++ Standard Library
- ▶ Above methods generate random numbers sequentially.

# Random Number Generation in MKL

- ▶ Intel MKL (Math Kernel Library)<sup>7</sup> provides efficient (pseudo) random number generators.
- ▶ Generates random numbers in serial/parallel mode.
- ▶ Generates a sequence of random numbers using one call.
- ▶ To generate random numbers using MKL:
  1. Create and initialize a stream
  2. Generate random numbers
  3. Delete the stream
- ▶ Related functions are defined in `<mk1_vsl.h>`

---

<sup>7</sup><https://www.intel.com/content/www/us/en/developer/tools/oneapi/onemkl.html#gs.aqzyf>

# RNG: Example 1

- ▶ Create a stream:

```
VSLStreamStatePtr stream;
```

- ▶ Initializer:

- ▶ Takes the stream, a uniform random number generator, and a seed as arguments.
- ▶ The seed initializes the generator and allows us to generate the same, different or unique sequences.

```
int seed = 777;
vslNewStream(&stream, VSL_BRNG_MT19937, seed);
```

We are using the Mersenne Twister pseudo-random generator above.

- ▶ Full list of supported generators:

```
https://www.intel.com/content/www/us/en/docs/onemkl/
developer-reference-c/2024-1/basic-generators.html
```

- ▶ This example shows how to generate random numbers from the standard normal distribution:<sup>8</sup>
  - ▶ Here we use ICDF (Inverse Cumulative Distribution Function)
  - ▶ Generate 1000 random numbers from  $N(0, 1)$  and store them in an (rands) array

```
float mean = 0;
float stdev = 1.0f;

float* rands = new float[1000];
vsRngGaussian(VSL_RNG_METHOD_GAUSSIAN_ICDF,
 stream,
 1000, rands,
 mean, stdev);
```

---

<sup>8</sup><https://www.intel.com/content/www/us/en/docs/onemkl/developer-reference-c/2024-1/random-number-generators-naming-conventions.html#TBL10-1>



- ▶ Use the random numbers:

```
for (int i=0; i<10; ++i)
{
 cout << rand[i] << endl;
}
```

- ▶ Delete the stream:

```
vslDeleteStream(&stream);
```

# MKL: Distributions

- ▶ MKL supports several distributions:  
`https://www.intel.com/content/www/us/en/docs/  
onemkl/developer-reference-c/2024-1/  
distribution-generators.html`


## Parallel RNG

# Parallel RNG

- ▶ Parallel simulations require us to use several streams.
  - ▶ E.g. each task/thread can use own parallel stream.
- ▶ Each parallel stream has to generate (pseudo) random numbers from the same underlying distribution.
  - ▶ I.e. each stream has to generate a non overlapping sub-sequence from the same underlying distribution.
- ▶ We can create such streams for parallel use, by:
  - ▶ Block splitting
  - ▶ Leap frogging


# Block Splitting

- ▶ Suppose the original sequence generates:  $x_1, x_2, x_3, x_n$  etc.
- ▶ One way to split an original sequence into non overlapping sub-sequences is:
  - ▶ 1<sup>st</sup> stream generates:  $x_1, x_2, x_k$
  - ▶ 2<sup>nd</sup> stream generates:  $x_{k+1}, x_{k+2}, x_{2k}$
  - ▶ 3<sup>rd</sup> stream generates:  $x_{2k+1}, x_{2k+2}, x_{3k}$  etc.
- ▶ Block splitting is to split a sequence into  $k$  non overlapping blocks, each generated by different processors.
- ▶ Example below shows how to split a sequence into 3 non overlapping streams using block-splitting.



# Leap Frogging

- ▶ Suppose the original sequence generates:  $x_1, x_2, x_3, x_n$  etc.
- ▶ Another way to split it into  $k$  non overlapping streams is to use how we deal a deck of cards to card players.
  - ▶ 1<sup>st</sup> stream generates:  $x_1, x_{k+1}, x_{2k+1}$
  - ▶ 2<sup>nd</sup> stream generates:  $x_2, x_{k+2}, x_{2k+2}$
  - ▶ 3<sup>rd</sup> stream generates:  $x_3, x_{k+3}, x_{2k+3}$  etc.
- ▶ This method is known as leap frogging.
- ▶ Example below shows how to split a sequence into 3 non overlapping streams using leap-frogging.



► Actual example:

```
[chanaka@midway2-login1 ParallelMC]$./rands_demo
Original Random Number Sequence
-4.957175, 0.693841, -1.582070, 1.056246, 2.182019, 0.823269, 0.619623, -1.837416, -0.497949, -0.107714, 0.909523, 0.756416, 1.193423, 0.383230, 1.901276, -0.501114, -0.884327, 0.051597, -0.273206, -0.882753,

Block Splitting
Seq 1: -4.957175, 0.693841, -1.582070, 1.056246, 2.182019,
Seq 2: 0.823269, 0.619623, -1.837416, -0.497949, -0.107714,
Seq 3: 0.909523, 0.756416, 1.193423, 0.383230, 1.901276,

Leap Frogging Sub Sequence
Seq 1: -4.957175, 1.056246, 0.619623, -0.107714, 1.193423, -0.501114, -0.273206, -0.271924, -1.232288, 0.376875,
Seq 2: 0.693841, 2.182019, -1.837416, 0.909523, 0.383230, -0.884327, -0.882753, -0.513166, 1.839407, -0.305231,
Seq 3: -1.582070, 0.823269, -0.497949, 0.756416, 1.901276, 0.051597, -0.098049, 0.371569, 0.141641, 0.723966,
[chanaka@midway2-login1 ParallelMC]$
```

► See demo for details.

# Using MKL and TBB on Midway

We need load intel and mkl modules first:

```
module use /software/intel/oneapi_hpc_2022.1/modulefiles
module load intel/2022.0
module load mkl/latest
```

- ▶ To enable MKL:

```
-qmkl[=<arg>]
 link to the Intel Intel MKL and bring
 in the associated headers
 parallel - link using the threaded Intel
 MKL libraries. This is the default
 when -qmkl is specified
 sequential - link using the non-threaded
 Intel MKL libraries
```

```
icc -qmkl rand_demo.cpp -o rand_demo
```

- ▶ Additionally, if you use TBB, to load TBB:

```
module load tbb/latest
```

To compile using tbb:

```
icc -qtbb
```




# Block Splitting vs Leap Frogging

- ▶ "Not all BRNGs support both these methods of generating independent sub-sequences. "
  - `https://cdrdv2-public.intel.com/671322/vslnotes-2.pdf`

## Example: Monte Carlo $\pi$

- ▶ You wrote a program (Assignment D) to estimate  $\pi$  using Monte Carlo by:
  - ▶ The figure below shows a circle with radius  $r = 1$  inscribed within a square.



- ▶ Pick  $N$  random points (given by two random  $x, y$  coordinates) within the square.
- ▶ Count the number of points inside the circle.

- ▶ Code below uses parallel random number streams to estimate the value of  $\pi$  using Monte Carlo:
  - ▶ Split the simulation into tasks
  - ▶ Each task uses a parallel stream to generate random numbers
  - ▶ A task is implemented as a function object
  - ▶ Use TBB task\_group to manage parallel tasks

```
tbb::task_group group;
long samples_per_task = total_samples / num_tasks;
long points_inside_circle_per_task[num_tasks];

for (int i = 0; i < num_tasks; i++)
{
 group.run(PiTask(samples_per_task,
 stream[i],
 points_inside_circle_per_task[i]));
}

group.wait();
```

- ▶ See demo for details. To build:  
`icc -qmk1 -qtbb mc_pi.cpp -o mc_pi`

# CPU Utilization

## ► Program uses more than one CPU core now:

### ⌵ CPU Usage Histogram 📄

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the total.

